


Machine Translation



Many slides from John Dehiero and Philip Koehn

Translation Task

- Text is both the input and the output.
- Input and output have roughly the same information content.
- Output is more predictable than a language modeling task.
- Lots of naturally occurring examples (but not much metadata).

Translation Examples

English-German News Test 2013 (a standard dev set)

Republican leaders justified their policy by the need to combat electoral fraud.

Die Führungskräfte der Republikaner
The Executives of the republican
rechtfertigen ihre Politik mit der
justify your politics With of the
Notwendigkeit , den Wahlbetrug zu
need , the election fraud to
bekämpfen .
fight .

Variety in Translations?

Human-generated reference translation

A small planet, whose is as big as could destroy a middle sized city, passed by the earth with a distance of 463 thousand kilometers. This was not found in advance. The astronomers got to know this incident 4 days later. This small planet is 50m in diameter. The astonomists are hard to find it for it comes from the direction of sun.

A commercial system from 2002

A volume enough to destroy a medium city small planet is big, flit earth within 463,000 kilometres of close however were not in advance discovered, astronomer just knew this matter after four days. This small planet diameter is about 50 metre, from the direction at sun, therefore astronomer very hard to discovers it.

Google Translate, 2020

An asteroid that was large enough to destroy a medium-sized city, swept across the earth at a short distance of 463,000 kilometers, but was not detected early. Astronomers learned about it four days later. The asteroid is about 50 meters in diameter and comes from the direction of the sun, making it difficult for astronomers to spot it.

Evaluation

BLEU Score

BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty (harshly penalizes translations shorter than the reference).

$$\text{Matched}_i = \sum_t \min \left\{ C_h(t_i), \max_j C_j(t_i) \right\}$$

If "of the" appears twice in hypothesis h but only at most once in a reference, then only the first is "correct"

$$P_i = \frac{\text{Matched}_i}{H_i}$$

"Clipped" precision of n-gram tokens

$$B = \exp \left\{ \min \left(0, \frac{n - L}{n} \right) \right\}$$

Brevity penalty only matters if the hypothesis corpus is shorter than the sum of (shortest) references.

$$\text{BLUE} = B \left(\prod_{i=1}^4 P_i \right)^{\frac{1}{4}}$$

BLEU is a mean of clipped precisions, scaled down by the brevity penalty.

Evaluation with BLEU

In this sense, the measures will partially undermine the American democratic system.

In this sense, these measures partially undermine the democratic system of the United States.




BLEU = 26.52, 75.0/40.0/21.4/7.7 (BP=1.000, ratio=1.143, hyp_len=16, ref_len=14)

(Papineni et al., 2002) BLEU is a method for automatic evaluation of machine translation.

Corpus BLEU Correlations with Average Human Judgments

There are ecological correlations over multiple segments; segment-level BLEU scores are noisy.

Commercial machine translation providers seem to all perform human evaluations of some sort.



(Ma et al., 2019) Results of the WMT19 Metrics Shared Task: Segment-Level and Strong MT Systems Pose Big Challenges

Human Evaluations

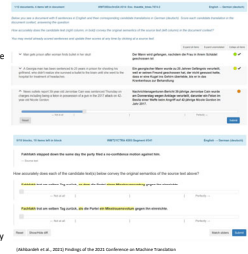
Direct assessment: adequacy & fluency

- Monolingual: Ask humans to compare machine translation to a human-generated reference. (Easier to source annotators)
- Bilingual: Ask humans to compare machine translation to the source sentence that was translated. (Compares to human quality)
- Annotators can assess segments (sentences) or whole documents.
- Segments can be assessed with or without document context.

Ranking assessment:

- Raters are presented with 2 or more translations.
- A human-generated reference may be provided, along with the source.
- "In a pairwise ranking experiment, human raters assessing adequacy and fluency show a stronger preference for human over machine translation when evaluating documents as compared to isolated sentences." (Laubli et al., 2018)

Editing assessment: How many edits required to reach human quality



Translationese and Evaluation

Translated text can: (Baker et al., 1993; Graham et al., 2019)

- be more explicit than the original source
- be less ambiguous
- be simplified (lexical, syntactically and stylistically)
- display a preference for conventional grammaticality
- avoid repetition
- exaggerate target language features
- display features of the source language

"If we consider only original source text (i.e. not translated from another language, or translationese), then we find evidence showing that human parity has not been achieved." (Torral et al., 2018)

(Baker et al., 1993) Target register and transfer: from quality implications and applications. (Baker et al., 1993) Translationese and machine translation evaluation. (Torral et al., 2018) Measuring the unlikelihood? Revisiting Claims of Human Parity in Human Machine Translation.

How are We Doing? Example: WMT 2019 Evaluation

2019 segment-in-context direct assessment (Barraut et al., 2019):

- ✓ German to English: many systems are tied with human performance;
- ✗ English to Chinese: all systems are outperformed by the human translator;
- ✗ English to Czech: all systems are outperformed by the human translator;
- ✗ English to Finnish: all systems are outperformed by the human translator;
- ✓ English to German: Facebook-FAIR achieves super-human translation performance; several systems are tied with human performance;
- ✗ English to Gujarati: all systems are outperformed by the human translator;
- ✗ English to Kazakh: all systems are outperformed by the human translator;
- ✗ English to Lithuanian: all systems are outperformed by the human translator;
- ✓ English to Russian: Facebook-FAIR is tied with human performance.


Statistical Machine Translation
(1990 - 2015)



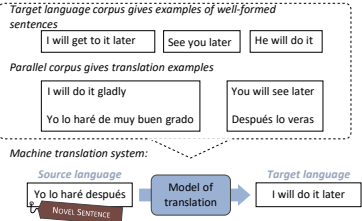
When I look at an article in Russian, I say:
"This is really written in English, but it has
been coded in some strange symbols. I
will now proceed to decode."

Warren Weaver (1949)


Levels of Transfer: Vauquois Triangle (1968)



Data-Driven Machine Translation



Stitching Together Fragments



Evolution of the Noisy Channel Model

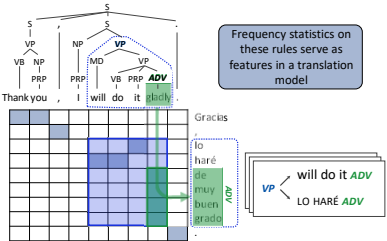
$$P(e|f) \propto P(f|e) \cdot P(e)$$
$$P(e|f) \propto P(f|e)^{\phi_{em}} \cdot P(e)^{\phi_{im}}$$
$$P(e|f) \propto \exp \left\{ \sum_t w_t \cdot f_t(e, f) \right\}$$

Chosen to minimize loss

E.g., log P(e)

Word Alignment and Phrase Extraction


Extracting Translation Rules



Counting Aligned Phrases

d'assister à la reunion et ||| to attend the meeting and
assister à la reunion ||| attend the meeting
la reunion et ||| the meeting and
nous ||| we
...

- Relative frequencies are the most important features in a phrase-based or syntax-based model.
- Scoring a phrase under a lexical model is the second most important feature.
- Estimation does not involve choosing among segmentations of a sentence into phrases.




Translation Options



- Many translation options to choose from
 - in Europarl phrase table: 2727 matching phrase pairs for this sentence
 - by pruning to the top 20 per phrase, 202 translation options remain


Decoding: Find Best Path



Phrase-Based Decoding

这	7人	中	包括	来自	法国	和	俄罗斯	的	宇航	员
1	2	3	4	5	6	7	8	9	10	11
the	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
2	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
3	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
4	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
5	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
6	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
7	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
8	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
9	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
10	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of
11	7 people	including	by	from	and	the	cosmonauts	the	astronauts	of

Machine Translation




Many slides from John Dehiero and Philip Koehn


Word Alignments

Word Alignment

Given a sentence pair, which words correspond to each other?




Word Alignment?



Is the English word **does** aligned to the German **wohnt** (verb) or **nicht** (negation) or neither?

Word Alignment?



How do the idioms **kicked the bucket** and **biss ins grass** match up?
Outside this exceptional context, **bucket** is never a good translation for **grass**

Lexical Translation / Word Alignment Models

Unsupervised Word Alignment

- Input: a *bitext*: pairs of translated sentences

nous acceptons votre opinion .
we accept your view .

- Output: *alignments*: pairs of translated words
- When words have unique sources, can represent as a (forward) alignment function a from French to English positions

Word Alignment

- Even today models are often built on the IBM alignment models
- Create probabilistic word-level translation models
- The models incorporate latent (unobserved) word alignments
- Optimize the probability of the observed words
- Use the imputed alignments to reveal word-level correspondence
- Throw out the translation models themselves

Alignment

- In a parallel text (or when we translate), we align words in one language with the words in the other

- Word positions are numbered 1-4

Alignment Function

- Formalizing alignment with an alignment function
- Mapping an English target word at position i to a German source word at position j with a function $a: i \rightarrow j$
- Example
 $a: \{1 \rightarrow 1, 2 \rightarrow 2, 3 \rightarrow 3, 4 \rightarrow 4\}$

Reordering

Words may be reordered during translation

$a: \{1 \rightarrow 3, 2 \rightarrow 4, 3 \rightarrow 2, 4 \rightarrow 1\}$


One-to-Many Translation

A source word may translate into multiple target words

$a: \{1 \rightarrow 1, 2 \rightarrow 2, 3 \rightarrow 3, 4 \rightarrow 4, 5 \rightarrow 4\}$

Dropping Words


Words may be dropped when translated
(German article *das* is dropped)



$a : \{1 \rightarrow 2, 2 \rightarrow 3, 3 \rightarrow 4\}$

Inserting Words

- Words may be added during translation
 - The English *just* does not have an equivalent in German
 - We still need to map it to something: special NULL token




$a : \{1 \rightarrow 1, 2 \rightarrow 2, 3 \rightarrow 3, 4 \rightarrow 0, 5 \rightarrow 4\}$

IBM Model 1: Allocation

IBM Model 1 (Brown 93)

- Alignments: a hidden vector called an *alignment* specifies which English source is responsible for each French target word.



Example

das		Haus		ist		klein	
e	t(e f)	e	t(e f)	e	t(e f)	e	t(e f)
the	0.7	house	0.8	is	0.8	small	0.4
that	0.15	building	0.16	's	0.16	little	0.4
which	0.075	home	0.02	exists	0.02	short	0.1
who	0.05	household	0.015	has	0.015	minor	0.06
this	0.025	shell	0.005	are	0.005	petty	0.04

$$p(e, a|f) = \frac{e}{4^3} \times t(\text{the}|\text{das}) \times t(\text{house}|\text{Haus}) \times t(\text{is}|\text{ist}) \times t(\text{small}|\text{klein})$$

$$= \frac{e}{4^3} \times 0.7 \times 0.8 \times 0.8 \times 0.4$$

$$= 0.0028e$$

IBM Models 1/2



Model Parameters


Translation: $P(F_1 = \text{Gracias} | E_{A_1} = \text{Thank})$ Alignment: $P(A_2 = 3)$

Expectation Maximization

EM Algorithm

- Incomplete data
 - if we had *complete data*, would could estimate *model*
 - if we had *model*, we could fill in the *gaps in the data*
- Expectation Maximization (EM) in a nutshell
 1. initialize model parameters (e.g. uniform)
 2. assign probabilities to the missing data
 3. estimate model parameters from completed data
 4. iterate steps 2-3 until convergence

EM Algorithm



- Initial step: all alignments equally likely
- Model learns that, e.g., *la* is often aligned with *the*

EM Algorithm



- After one iteration
- Alignments, e.g., between *la* and *the* are more likely

EM Algorithm



- After another iteration
- It becomes apparent that alignments, e.g., between *fleur* and *flower* are more likely (pigeon hole principle)

EM Algorithm



- Convergence
- Inherent hidden structure revealed by EM


... la maison ... la maison bleu ... la fleur ...
 // | X |
 ... the house ... the blue house ... the flower ...

```
p(la|the) = 0.453
p(le|the) = 0.334
p(maison|house) = 0.876
p(bleu|blue) = 0.563
...
```

- | | initial | 1st it. | 2nd it. | 3rd it. | ... | final |
|------------------------------|---------|---------|---------|---------|-----|--------|
| <i>p</i> (the haus/das haus) | 0.0625 | 0.1875 | 0.1905 | 0.1913 | ... | 0.1875 |
| <i>p</i> (the book/das buch) | 0.0625 | 0.1406 | 0.1790 | 0.2075 | ... | 0.25 |
| <i>p</i> (a book/ein buch) | 0.0625 | 0.1875 | 0.1907 | 0.1913 | ... | 0.1875 |
| perplexity | 4092 | 293 | 153.6 | 131.6 | ... | 113.8 |


Problems with Model 1

- There's a reason they designed models 2-5!
- Problems: alignments jump around, align everything to rare words
- Experimental setup:
 - Training data: 1.1M sentences of French-English text, Canadian Hansards
 - Evaluation metric: alignment error Rate (AER)
 - Evaluation data: 447 hand-aligned sentences




IBM Model 2: Global Monotonicity

Monotonic Translation



Local Order Change



IBM Model 2

- Alignments tend to the diagonal (broadly at least)

$$P(f, a|e) = \prod_j P(a_j = i|j, I, J) P(f_j|e_i)$$

$$P(\text{dist} = i - j \frac{I}{J})$$


$$\frac{1}{Z} e^{-\alpha(i-j)\frac{I}{J}}$$

EM for Models 1/2

- Model 1 Parameters:
 - Translation probabilities (1+2) $P(f_j|e_i)$
 - Distortion parameters (2 only) $P(a_j = i|j, I, J)$
 - Start with $P(f_j|e_i)$ uniform, including $P(f_j|null)$
 - For each sentence:
 - For each French position j
 - Calculate posterior over English positions
- $$P(a_j = i|f, e) = \frac{P(a_j = i|j, I, J) P(f_j|e_i)}{\sum_i P(a_j = i|j, I, J) P(f_j|e_i)}$$
- (or just use best single alignment)
 - Increment count of word f_j with word e_i by these amounts
 - Also re-estimate distortion probabilities for model 2
- Iterate until convergence

HMM Model: Local Monotonicity

Phrase Movement



The HMM Model



Model Parameters

Emissions: $P(F_1 = \text{Gracias} | E_{A_1} = \text{Thank})$ Transitions: $P(A_2 = 3 | A_1 = 1)$

The HMM Model

- Model 2 preferred global monotonicity
- We want local monotonicity:
 - Most jumps are small
- HMM model (Vogel 96)

f	$t(f e)$
nationale	0.469
national	0.418
nationaux	0.054
nationales	0.029

$$P(f, a|e) = \prod_j P(a_j | a_{j-1}) P(f_j | e_j)$$


$$P(a_j - a_{j-1})$$



- Re-estimate using the forward-backward algorithm
- Handling nulls requires some care
- What are we still missing?

Models 3+: Fertility

IBM Models 3/4/5



[from Al-Onaizan and Knight, 1998]

Examples: Translation and Fertility

<i>the</i>			
<i>f</i>	$t(f e)$	ϕ	$n(\phi e)$
le	0.497	1	0.746
la	0.207	0	0.254
les	0.155		
l'	0.086		
ce	0.018		
cette	0.011		

<i>not</i>			
<i>f</i>	$t(f e)$	ϕ	$n(\phi e)$
ne	0.497	2	0.735
pas	0.442	0	0.154
non	0.029	1	0.107
rien	0.011		

farmers

<i>f</i>	$t(f e)$	ϕ	$n(\phi e)$
agriculteurs	0.442	2	0.731
les	0.418	1	0.228
cultivateurs	0.046	0	0.039
producteurs	0.021		

Example: Idioms

he is nodding
 ↓
 il hoche la tête

<i>nodding</i>			
<i>f</i>	$t(f e)$	ϕ	$n(\phi e)$
signe	0.164	4	0.342
la	0.123	3	0.293
tête	0.097	2	0.167
oui	0.086	1	0.163
fait	0.073	0	0.023
que	0.073		
hoche	0.054		
hocher	0.048		
faire	0.030		
me	0.024		
approuve	0.019		
qui	0.019		
un	0.012		
faites	0.011		

Example: Morphology

<i>should</i>			
<i>f</i>	$t(f e)$	ϕ	$n(\phi e)$
devrait	0.330	1	0.649
devraient	0.123	0	0.336
devrions	0.109	2	0.014
faudrait	0.073		
faut	0.058		
doit	0.058		
aurait	0.041		
doivent	0.024		
devons	0.017		
devrais	0.013		

Getting Phrases


Word Alignment with IBM Models

- IBM Models create a **many-to-one** mapping
 - words are aligned using an alignment function
 - a function may return the same value for different input (one-to-many mapping)
 - a function can not return multiple values for one input (no many-to-one mapping)
- Real word alignments have **many-to-many** mappings


Symmetrization

- Run IBM Model training in both directions
 - two sets of word alignment points
- Intersection: high precision alignment points
- Union: high recall alignment points
- Refinement methods explore the sets between intersection and union

Example




Growing Heuristics



- black: intersection grey: additional points in union
- Add alignment points from union based on heuristics:
 - directly/diagonally neighboring points
 - finally, add alignments that connect unaligned words in source and/or target
- Popular method: grow-diag-final-and

Phrase-Based Model



- Foreign input is segmented in phrases
- Each phrase is translated into English
- Phrases are reordered

Phrase Translation Table

- Main knowledge source: table with phrase translations and their probabilities
- Example: phrase translations for *natuerlich*

Translation	Probability $\phi(e f)$
of course	0.5
naturally	0.3
of course ,	0.15
, of course ,	0.05

Scoring Phrase Translations

- Phrase pair extraction: collect all phrase pairs from the data
- Phrase pair scoring: assign probabilities to phrase translations
- Score by relative frequency:

$$\phi(f|e) = \frac{\text{count}(e, f)}{\sum_i \text{count}(e, f_i)}$$


Real Example

- Phrase translations for *den Vorschlag* learned from the Europarl corpus:

English	$\phi(e f)$	English	$\phi(e f)$
the proposal	0.6227	the suggestions	0.0114
's proposal	0.1068	the proposed	0.0114
a proposal	0.0341	the motion	0.0091
the idea	0.0250	the idea of	0.0091
this proposal	0.0227	the proposal ,	0.0068
proposal	0.0205	its proposal	0.0068
of the proposal	0.0159	it	0.0068
the proposals	0.0159


- lexical variation (proposal vs suggestions)
- morphological variation (proposal vs proposals)
- included function words (the, a, ...)
- noise (it)

Extracting Phrase Pairs




extract phrase pair consistent with word alignment:
assumes that / geht davon aus , dass

Consistent



All words of the phrase pair have to align to each other.


Phrase Pair Extraction



Smallest phrase pairs:
michael — michael
assumes — geht davon aus / geht davon aus ,
that — dass / , dass
he — er
will stay — bleibt
in the — im
house — haus

unaligned words (here: German comma) lead to multiple translations

Larger Phrase Pairs



michael assumes — michael geht davon aus / michael geht davon aus ,
assumes that — geht davon aus , dass ; assumes that he — geht davon aus , dass er
that he — dass er / , dass er ; in the house — im haus
michael assumes that — michael geht davon aus , dass
michael assumes that he — michael geht davon aus , dass er
michael assumes that he will stay in the house — michael geht davon aus , dass er im haus bleibt
assumes that he will stay in the house — geht davon aus , dass er im haus bleibt
that he will stay in the house — dass er im haus bleibt ; dass er im haus bleibt ,
he will stay in the house — er im haus bleibt ; will stay in the house — im haus bleibt


Other Scoring Terms

More Feature Functions

- Bidirectional alignment probabilities: $\phi(e|f)$ and $\phi(f|e)$
- Rare phrase pairs have unreliable phrase translation probability estimates
→ lexical weighting with word translation probabilities


$$\text{lex}(e|f, a) = \prod_{i=1}^{\text{length}(a)} \frac{1}{| \{ j | (i, j) \in a \} |} \sum_{v(i, j) \in a} w(e_i | f_j)$$

Decoding: Hypothesis Expansion




pick any translation option, create new hypothesis

Decoding: Hypothesis Expansion




create hypotheses for all other translation options

Decoding: Hypothesis Expansion



also create hypotheses from created partial hypothesis

Decoding: Find Best Path



backtrack from highest scoring complete hypothesis


Dynamic Programming

Computational Complexity


- The suggested process creates exponential number of hypothesis
- Machine translation decoding is NP-complete
- Reduction of search space:
 - recombination (risk-free)
 - pruning (risky)

Recombination


- Two hypothesis paths lead to two matching hypotheses
 - same foreign words translated
 - same English words in the output



- Worse hypothesis is dropped



Stacks



- Hypothesis expansion in a stack decoder
 - translation option is applied to hypothesis
 - new hypothesis is dropped into a stack further down

Stack Decoding Algorithm

```

1: place empty hypothesis into stack 0
2: for all stacks (0..n - 1) do
3:   for all hypotheses in stack do
4:     for all translation options do
5:       if applicable then
6:         create new hypothesis
7:         place in stack
8:         recombine with existing hypothesis if possible
9:         prune stack if too big
10:      end if
11:    end for
12:  end for
13: end for

```

Pruning

Pruning

- Pruning strategies
 - histogram pruning: keep at most k hypotheses in each stack
 - stack pruning: keep hypothesis with score $\alpha \times$ best score ($\alpha < 1$)
- Computational time complexity of decoding with histogram pruning


$$O(\text{max stack size} \times \text{translation options} \times \text{sentence length})$$
- Number of translation options is linear with sentence length, hence:

$$O(\text{max stack size} \times \text{sentence length}^2)$$
- Quadratic complexity

Future Costs

Translating the Easy Part First?

the tourism initiative addresses this for the first time




both hypotheses translate 3 words
worse hypothesis has better score

Estimating Future Cost

- Future cost estimate: how expensive is translation of rest of sentence?
- Optimistic: choose cheapest translation options
- Cost for each translation option
 - translation model: cost known
 - language model: output words known, but not context
→ estimate without context
 - reordering model: unknown, ignored for future cost estimation

Cost Estimates from Translation Options



cost of cheapest translation options for each input span (log-probabilities)


Cost Estimates for all Spans

- Compute cost estimate for all contiguous spans by combining cheapest options

first word	1	2	3	4	5	6	7	8	9
the	-1.0	-3.0	-4.5	-6.9	-8.3	-9.3	-9.6	-10.6	-10.6
tourism	-2.0	-3.5	-5.9	-7.3	-8.3	-8.6	-9.6	-9.6	
initiative	-1.5	-3.9	-5.3	-6.3	-6.6	-7.6	-7.6		
addresses	-2.4	-3.8	-4.8	-5.1	-6.1				
this	-1.4	-2.4	-2.7	-3.7	-3.7				
for	-1.0	-1.3	-2.3	-2.3					
the	-1.0	-2.2	-2.3						
first	-1.9	-2.4							
time	-1.6								


- Function words cheaper (the: -1.0) than content words (tourism: -2.0)
- Common phrases cheaper (for the first time: -2.3) than unusual ones (tourism initiative addresses: -5.9)

Combining Score and Future Cost



- Hypothesis score and future cost estimate are combined for pruning
 - left hypothesis starts with hard part: **the tourism initiative**
score: -5.88, future cost: -6.1 → total cost -11.98
 - middle hypothesis starts with easiest part: **this for ... time**
score: -4.11, future cost: -9.3 → total cost -13.41
 - right hypothesis picks easy parts: **this for ... time**
score: -4.86, future cost: -9.1 → total cost -13.96

A* Search



- Uses *admissible* future cost heuristic: never overestimates cost
- Translation agenda: create hypothesis with lowest score + heuristic cost
- Done, when complete hypothesis created